

# 535 Project Final Report

Shankar Subramanian and Stelios Papoutsakis

## Architecture Overview:

Our architecture is a stack based machine with emphasis on crypto math instructions. It is developed as a clean slate architecture. Our inspiration is from the Bitcoin/Ethereum stack machines.

### Word size:

32 bits

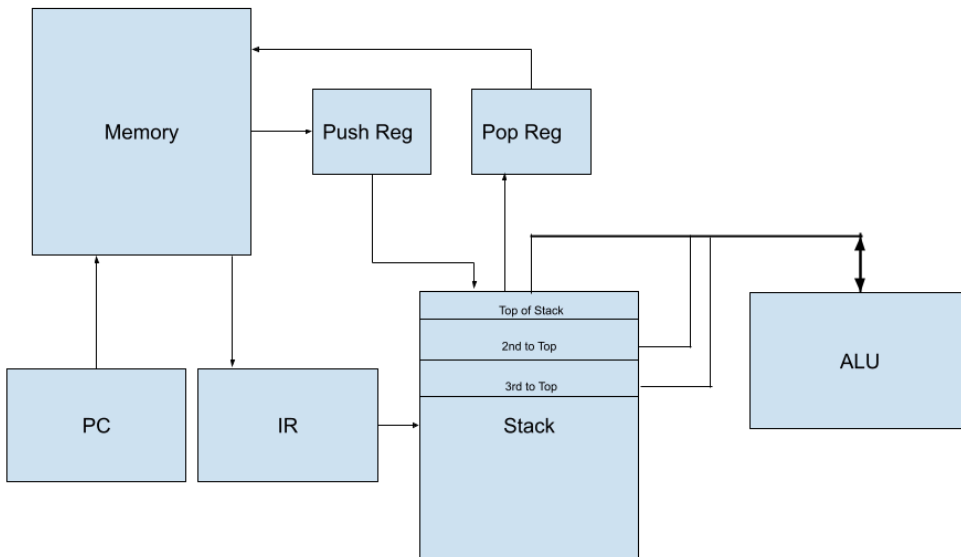
### Instruction size:

8 bits

### Supported data types:

1. Unsigned Integer ( 32 bit )
2. Block for AES ( 128 bit )
3. Block for RSA ( 64 bit )
4. Group Element ( 64 bit )

## Stack Machine Organization



As shown above, the ALU is able to operate on the top three stack elements. It then writes its computation directly to the top of the stack. The push and pop registers are used to access and write values back to memory. If the stack overflows, the oldest element inserted into the stack ( element 32 ) is overwritten.

## Special Registers:

1. Program Counter
2. Offset Register
3. 32x 128 bit stack registers as a register file
4. Pop Register ( 128 bit )
5. Push Register ( 128 bit )
6. Link Register

There are no general purpose registers in our design.

The offset register contains the byte offset ( high to low ) for an instruction in the word pointed to by PC

The execution model is a Stack Based Machine, fetching four instructions per word.

We will be using the Princeton memory organization.

Our address space would be 65532 ( $2^{16}$ ) words, though not all of it will be used in simulation.

We will be doing word addressing.

We will use indirect addressing for all instructions, with the address being read from the stack.

## Pipeline:

Our architecture features a 3 stage pipeline, described below

*Fetch*: Fetches instruction based on PC and OFFSET. Issues a byte to decode

*Decode*: Decodes the byte issued from fetch. Pops off operands from the stack if necessary and classifies the instructions as “Alu”, “Branch”, “Memory” or “Crypto”

*Execute*: Executes the instruction from Decode. Based on the classification of the instruction it feeds it to the appropriate executor module. It uses a lookup table to identify how long certain non memory instructions should block for ( memory instructions block based on the status of memory ). Pushes any results generated immediately to the stack.

We opted for a 3 stage pipeline because of the high hazards a 5 stage one would provide. This is because most operations always write to the top of the stack, which would result in the top of stack being a pending register. This would result in the decode stage stalling often as the decode stage is responsible for popping operands from the top of stack. Modifying execute to write to the top of stack removed this hazard

## Cache:

We used a direct mapped write through no allocate cache. We only have one level of cache. The cache size is 16 lines or in other words 64 bytes.

## Instructions:

There are two types of instructions, ones that operate with immediate values and ones that exclusively use the stack

Since a majority of our instructions use the top stack values as input and push a new value on completion, there is no need to define multiple type fields that specify different operand formats. The amount of pops used for operands are defined in the implementation of the instruction itself.

Any value used by an instruction ( specified by stack Input: ) is removed from the stack as part of its execution. Additionally, instructions will use however many bits they need starting from the least significant bits. For example, unsigned integer arithmetic will only use 32 out of the 128

bits a stack register has as an operand value. Output lengths are specified in the sub sections below. Any output less than 128 bits is extended with 0s.

## Instructions With Immediates ( Type 0 ):

Format:

	7	6	5	4	3	2	1	0
Num	TYPE	OPCODE			Stack Position/Immediate			
0	0	PUSH_VAL			INTEGER_LSB			
1	0	SWAP			INTEGER_LSB			
2	0	DUP			INTEGER_LSB			

Opcode: 0

Mnemonic : PUSH\_VAL

Stack Input:

Stack Output: Immediate

Operation: Pushes the immediate value on top of the stack. The upper 122 significant bits are extended with 0s

Opcode: 1

Mnemonic : SWAP

Stack Input:

Stack Output:

Operation: Takes the value in the stack index by immediate, removes it and pushes it to the top of stack

Opcode: 2

Mnemonic : DUP

Stack Input:

Stack Output:

Operation: Takes the value in the stack index by immediate, duplicates it and pushes it on top of the stack

## Instructions OPCODE Only ( Type 1 ):

Format:



	7	6	5	4	3	2	1	0
Num	TYPE	OPCODE						
13	1	ADD						
14	1	MUL						
15	1	SUB						
16	1	DIV						
17	1	MOD						
18	1	AND						
19	1	OR						
20	1	XOR						
21	1	NOT						
22	1	EQ						
23	1	EQ 0						
24	1	GEQ						
25	1	LEQ						
26	1	GT						
27	1	LT						
28	1	R SHIFT						
29	1	L SHIFT						

	7	6	5	4	3	2	1	0
Num	TYPE	OPCODE						
47	1	NOOP						
48	1	HALT						
50	1	DUP_TOP						
51	1	SWAP TOP						

	7	6	5	4	3	2	1	0
Num	TYPE	OPCODE						
30	1	GCD						
31	1	LCM						
32	1	RAND_INT						
33	1	RAND_GRP						
34	1	RABIN						
35	1	MOD_ADD						
36	1	MOD_MUL						
37	1	MOD_INV						
38	1	DH						
39	1	RSA						
40	1	AESE						
41	1	AESD						
42	1	RAND_BLK						
43	1	XOR_BLK						
44	1	SCHNORR_SIG						
45	1	SCHNORR_VER						
46	1	SHA256						

Load/Store instructions:

Opcode: 0

Mnemonic : LDR\_32

Stack Input: Address

Stack Output: None

Operation: Loads a word from memory into the PUSH\_REGISTER register at the address on the top of the stack.

Opcode: 1

Mnemonic : STR\_32

Stack Input: Address

Stack Output: None

Operation : Stores the 32 LSB in the POP\_REGISTER at the address on the top of the stack.

Opcode: 2

Mnemonic : LDR\_64

Stack Input: Address

Stack Output: None

Operation: Loads two consecutive words from memory into the PUSH\_REGISTER register at the address on the top of the stack.

Opcode: 3

Mnemonic : STR\_64

Stack Input: Address

Stack Output: None

Operation: Stores the 64 LSB value in the POP\_REGISTER at the address on the top of the stack.

Opcode: 4

Mnemonic : LDR\_128

Stack Input: Address

Stack Output: None

Operation: Loads four consecutive words from memory into the PUSH\_REGISTER register at the address on the top of the stack.

Opcode: 5

Mnemonic : STR\_128

Stack Input: Address

Stack Output: None

Operation: Stores the 128 LSB value in the POP\_REGISTER at the address on the top of the stack.

Opcode: 6

Mnemonic : PUSH

Stack Input: None

Stack Output: Literal at top of the stack

Operation : Pushes the value located in the PUSH\_REGISTER to the top of the stack.

Opcode: 7

Mnemonic : POP

Stack Input: Value at the top of the stack

Stack Output: None

Operation: Pops the value at the top of the stack into the POP\_REGISTER



## Control instructions/modes:

Opcode: 8

Mnemonic : JMP

Stack Input: Address

Stack Output: None

Operation: An unconditional jump to the address at the top of the stack.

Opcode: 9

Mnemonic : JMP\_if\_1

Stack Input: condition Address

Stack Output: None

Operation: Tests to see if the top of stack's least significant bit is 1. Jumps to address stored at the second from the top of stack if true

Opcode: 10

Mnemonic : JMP\_if\_0

Stack Input: condition Address

Stack Output: None

Operation: Tests to see if the top of stack's least significant bit is 0. Jumps to address stored at the second from the top of stack if true

Opcode: 11

Mnemonic : SR

Stack Input: Address

Stack Output: None

Operation : Places current address in link register and sets the PC to address on the top of stack.

Opcode: 12

Mnemonic : RET

Stack Input: None

Stack Output: None

Operation: Replaces the PC with the value located in the link register

## Unsigned Integer ALU Instructions:

All operations result in 32 bit integer values. The 33rd bit is used as overflow if applicable. The exception is MULT which returns a 64 bit product.

Opcode: 13

Mnemonic: ADD

Stack Input: num1 num2

Stack Output: sum

Operation: Added the top two values on the stack and push their sum

Opcode: 14

Mnemonic: MUL

Stack Input: num1 num2

Stack Output: mult

Operation: Multiply the top two values on the stack and push a 64 bit product. Note that the stack uses 128 bit registers

Opcode: 15

Mnemonic: SUB

Stack Input: num1 num2

Stack Output: sub

Operation: Subtract the top two values on the stack and push their difference.

Opcode: 16

Mnemonic: DIV

Stack Input: num1 num2

Stack Output: quo

Operation: Divide the top two values on the stack and push the quotient. If the num2 is 0, this operation pushes zero onto the stack

Opcode: 17

Mnemonic: MOD

Stack Input: num mod

Stack Output: rem

Operation: Take the value from the 2nd top of stack to use as a mod value and use the top of stack as divisor. Pushes the result on top of the stack

Opcode: 18

Mnemonic: AND

Stack Input: num1 num2

Stack Output: res

Operation: Bitwise AND the top two values of the stack. Push the result on the top of the stack.

Opcode: 19

Mnemonic: OR

Stack Input: num1 num2

Stack Output: res

Operation: Bitwise OR the top two values of the stack. Push the result on the top of the stack.

Opcode: 20

Mnemonic: XOR

Stack Input: num1 num2

Stack Output: res

Operation: Bitwise XOR the top two values of the stack. Push the result on the top of the stack.

Opcode: 21

Mnemonic: NOT

Stack Input: num

Stack Output: res

Operation: Invert the top value of the stack.. Push the result on the top of the stack.

Opcode: 22

Mnemonic: EQ

Stack Input: num1 Num2

Stack Output: condition

Operation: Pushes 1 onto the stack if the top two values are equal, else pushes 0

Opcode: 23

Mnemonic: EQ\_0

Stack Input: num1

Stack Output: condition

Operation: Pushes 1 onto the stack if the top value equals 0, else pushes 0

Opcode: 24

Mnemonic: GEQ

Stack Input: num1 Num2

Stack Output: condition

Operation: Pushes 1 onto the stack if num1 is greater than or equal num2, else pushes 0

Opcode: 25

Mnemonic: LEQ

Stack Input: num1 Num2

Stack Output: condition

Operation: Pushes 1 onto the stack if num1 is less than or equal num 2, else pushes 0

Opcode: 26

Mnemonic: GT

Stack Input: num1 Num2

Stack Output: condition

Operation: Pushes 1 onto the stack if num1 is greater than num2, else pushes 0

Opcode: 27

Mnemonic: LT

Stack Input: num1 Num2

Stack Output: condition

Operation: Pushes 1 onto the stack if num1 is less than num2, else pushes 0

Opcode: 28

Mnemonic: R\_SHIFT

Stack Input: pad num

Stack Output: shifted\_num

Operation: Right shifts the element of the stack at depth 2 by the value at the top of the stack, appends 0 as pads to the left end.

Opcode: 29

Mnemonic: L\_SHIFT

Stack Input: pad num

Stack Output: shifted\_num

Operation: Left shifts the element of the stack at depth 2 by the value at the top of the stack, append 0 as pads to the right end.

Opcode: 30

Mnemonic: GCD

Stack Input: a b

Stack Output: gcd

Operation: Computes the GCD of the top 2 elements of the stack

Opcode: 31

Mnemonic: LCM

Stack Input: a b

Stack Output: lcm

Operation: Computes the LCM of the top 2 elements of the stack

Opcode: 47

Mnemonic: NOOP

Stack Input:

Stack Output:

Operation: Sends a NOOP through the pipeline

Opcode: 48

Mnemonic: HALT

Stack Input:

Stack Output:

Operation: Halts the pipeline's execution until run is triggered.

Opcode: 50

Mnemonic: DUP\_TOP

Stack Input: stack\_location

Stack Output: element at stack location

Operation: Computes the LCM of the top 2 elements of the stack

Opcode: 51

Mnemonic: SWAP\_TOP

Stack Input: stack\_location

Stack Output:

Operation: Takes the value in the stack index by immediate, and swaps it with the element at index 1 in the stack

## Crypto instructions:

Opcode: 32

Mnemonic: RAND\_INT

Stack Input: None

Stack Output: rand integer

Operation: Adds onto the stack a random 32 bit integer

Opcode: 33

Mnemonic: RAND\_GRP

Stack Input: None

Stack Output: rand group integer

Operation: Adds onto the stack a random 64 bit integer

Opcode: 34

Mnemonic: RABIN

Stack Input: a

Stack Output: 1 or 0

Operation: Checks if the number on top of the stack is a prime or not.

Opcode: 35

Mnemonic: MOD\_ADD

Stack Input: a b n

Stack Output: sum

Operation: Perform ADD and then modulus with the 3rd element in the stack.

Opcode: 36  
Mnemonic: MOD\_MUL  
Stack Input: a b n  
Stack Output: prod  
Operation: Performs MUL and then modulus with the 3rd element in the stack.

Opcode: 37  
Mnemonic: MOD\_INV  
Stack Input: a b n  
Stack Output: The inverse modulus as an unsigned integer  
Operation: Takes the top 3 elements of the stack as (group element, group element, modulo value) and computes the inverse modulus.

Opcode: 38  
Mnemonic: DH  
Stack Input: a b n  
Stack Output: key  
Operation: Takes in the top 3 elements of the stack as (group element, group element, modulo value) and computes the diffie-hellman key.

Opcode: 39  
Mnemonic: RSA  
Stack Input: m e n  
Stack Output: The RSA encryption/decryption as an unsigned integer  
Operation: Performs m raised to e and modulo in n

Opcode: 40  
Mnemonic: AESE  
Stack Input: m k  
Stack Output: c  
Operation: Computes the AES encryption of the top most value and uses the second from the top value as the key.

Opcode: 41  
Mnemonic: AESD  
Stack Input: c k  
Stack Output: m  
Operation: Computes the AES decryption on the top most value and uses the second from the top value as the key

Opcode: 42  
Mnemonic: RAND\_BLK  
Stack Input:

Stack Output: rand

Operation: Generates a block (128 bit) random integer

Opcode: 43

Mnemonic: XOR\_BLK

Stack Input: a b

Stack Output: m

Operation: Computes the XOR of the top 2 blocks on the stack

Opcode: 44

Mnemonic: SCHNORR\_SIG

Stack Input: m

Stack Output: sig

Operation: Computes the Schnorr signature of the topmost value of the stack with the key in the key register.

Opcode: 45

Mnemonic: SCHNORR\_VER

Stack Input: sig pub

Stack Output: 1 or 0

Operation: Verifies if the data at the top of the stack passes through the schnorr verification test with the public key as the 2nd element in the stack.

Opcode: 46

Mnemonic: SHA256

Stack Input: m

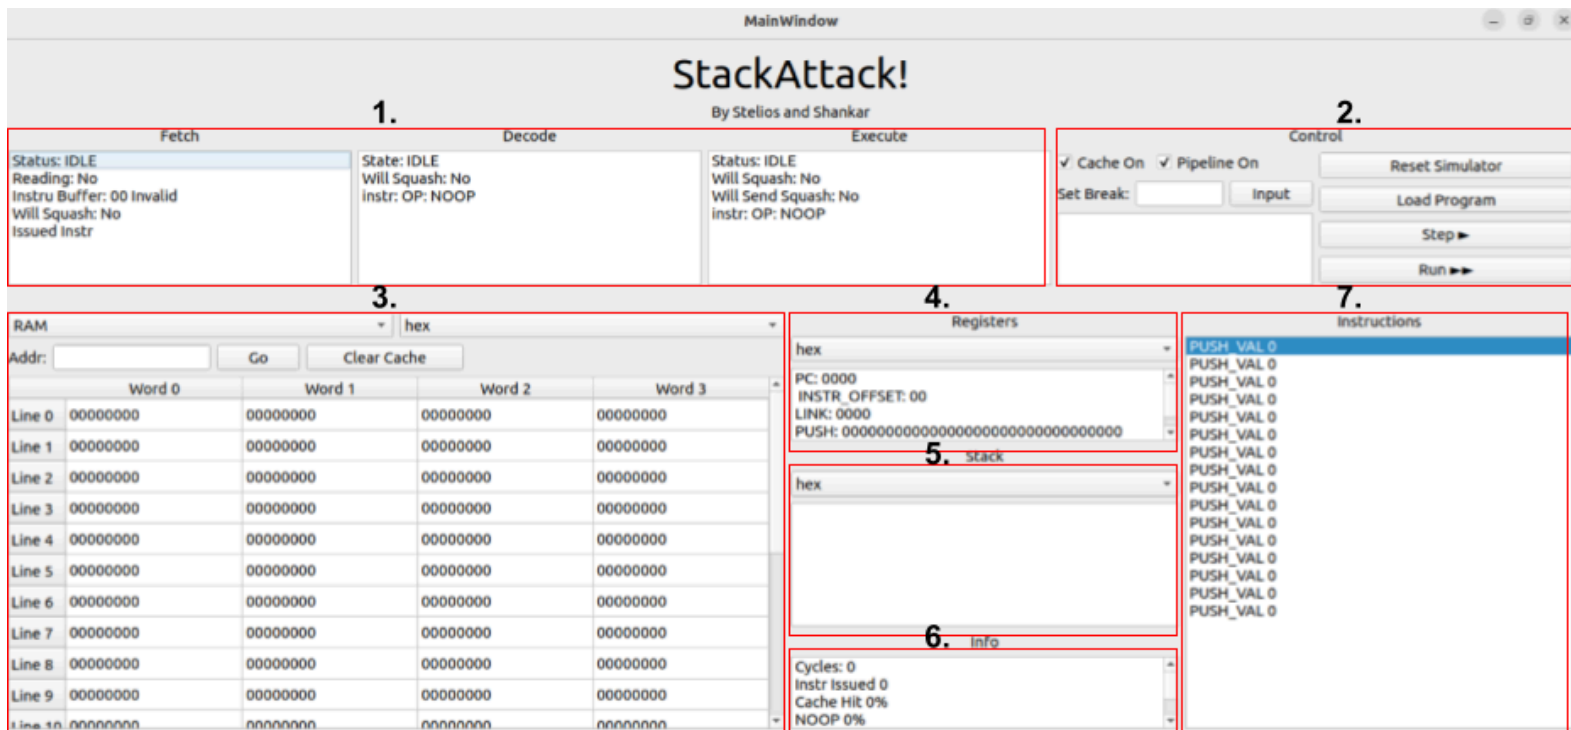
Stack Output: SHA256 hash

Operation: Hashes ( SHA256 ) the topmost element of the stack and puts it onto the stack

## Crypto Demo:

We plan on adding an ascii encoded mode in our GUI for the RAM + Cache. This will help us with our AES, RSA, schnorr signatures, and Hashing demo, which will show a simple message being encrypted, decrypted, hashed and signed.

# Simulator:



## Gui Description

Below is an explanation of the 7 main components

- 1. Pipeline State:** These windows show the internal state of each stage in the pipe. Information displayed includes status, squashing state, and instruction currently worked on
- 2. Control Panel:** This is where the user controls the simulation. This includes the ability to toggle the cache and pipeline, load a binary program, reset the simulator state, step through a program, or run a program. This is also where the user defines breakpoints in the format "ADDR, OFFSET", where ADDR is the address of the word containing the instruction and OFFSET is the byte within the word ( highest to lowest ). The user may double click on a displayed breakpoint to remove it
- 3. Memory Panel:** This displayed the memory. A user may type an address ( base 10 only ), into the search box and 20 lines starting at the requested address will be displayed. The user can switch between ram and cache change with the left combo box and change data representation with the right. The user may also clear the cache at any point
- 4. Register Panel:** Displays the register values. The user may change the data representation with the combo box
- 5. Stack Panel:** Displays the stack. The user may change the data representation with the combo box



6. Info Panel: Displays statistics about the simulation. This includes cycle count, instructions issued, Cache hit rate, and percent of time spent executing each instruction type
7. Disassembled Panel: Displays disassembled binary code. The code that is displayed is taken from the line that includes the address typed into the ADDR search box in 3. If the PC is pointing to a byte that is disassembled, the gui will highlight it. Any byte which is not a valid instruction is displayed as '-'

## User Manual

To run a program, a binary executable may be loaded using the “load” button in the control panel. To define breakpoints, the user may enter a value into the textbox formatted as “ADDR, OFFSET”, where ADDR is the address of the word containing the instruction and OFFSET is the byte within the word ( highest to lowest ). The user may then use “Step” to simulate a cycle, or “Run” to simulate the program until a breakpoint or HALT instruction is reached.

## SWE Methods:

### Tech Stack

The simulator was written exclusively in python. Qt for Python was used for a GUI library. Qt Designer was used as a drag and drop GUI creator. Pyuic5 was used to compile the XML produced by Qt Designer into Python.

### Version Control

Git and Github.

### Management Plan

Stelios was more focused on the GUI, Execution stage and writing benchmarks (AES, Exchange Sort, Standard Matrix Multiplication).

Shankar worked on designing the Fetch and Decode stages, testing, and a Matrix benchmark that tried to take advantage of the stack.

### Communication and Meeting Plan

Communication was done through discord. Communication occurs at least once a week which includes status updates and encountered difficulties. In person meetings were not regularly scheduled and only occurred in initial stages of pipeline design

## Division of Work:

Stelios:

- GUI Creation
- Execution Stage
- Memory
- AES, Exchange sort and standard matrix benchmark

Shankar:

- Fetch
- Decode
- Testing
- Stack Matrix Benchmark

## Benchmark Results:

Note that the cache delay was set to 3 cycles and main memory delay was set to 100 cycles for the benchmarks below

Matrix	Cache/Pipeline			
	On/On	On/Off	Off/On	Off/Off
Cycles	63,209.00	86,619.00	382,944.00	405,384.00
Inst Issued	12,173.00	11,873.00	12,173.00	11,873.00
Cache Hit	88.61%	88.58%	0.00%	0.00%
NOOP	33.81%	51.41%	80.28%	81.90%
ALU	17.30%	12.63%	2.86%	2.70%
BRANCH	0.51%	0.37%	0.08%	0.08%
CRYPTO	0.00%	0.00%	0.00%	0.00%
MEMORY	47.87%	34.90%	16.45%	15.17%

Matrices were of sizes:

- Matrix 1 - 20x5
- Matrix 2 - 3x20
- Resulting matrix - 3x5

The values were random positive integers generated in the range [0,100]

Insertion Sort	Cache/Pipeline			
	On/On	On/Off	Off/On	Off/Off
Cycles	1,280,274.00	1,648,448.00	7,755,889.00	7,185,444.00
Inst Issued	221,903.00	212,000.00	221,903.00	212,000.00
Cache Hit	93.36%	92.65%	0.00%	0.00%
NOOP	18.49%	39.81%	66.72%	83.70%
ALU	14.44%	11.21%	2.38%	2.57%
BRANCH	1.38%	1.07%	0.23%	0.25%
CRYPTO	0.00%	0.00%	0.00%	0.00%
MEMORY	61.46%	46.70%	17.91%	13.21%

The array used was of size 100 with random positive integers ranging from [0,1000]

AES (AESE + AESD)	Cache/Pipeline			
	On/On	On/Off	Off/On	Off/Off
Cycles	17,419.00	19,416.00	42,638.00	44,639.00
Instr Issued	1,049.00	1,020.00	1,049.00	1,020.00
Cache Hit	84.74%	84.74%	0.00%	0.00%
NOOP	13.94%	22.79%	61.79%	66.54%
ALU	5.05%	4.53%	2.06%	1.97%
BRANCH	0.17%	0.15%	0.07%	0.07%
CRYPTO	2.58%	2.31%	1.05%	1.01%
MEMORY	78.09%	69.92%	34.96%	30.29%

The message used was 1632-bits long and the key was 16-bits long.

The AES was run on CBC mode, which required a random IV (initialization vector) that was used to XOR the message block. This was generated using RAND\_INT.

The benchmark covered both encryption and decryption of the message

## Summary:

In conclusion, we learned that there are a lot of hazards that can impact the performance of our architecture. This is apparent with the high percentage of CPU idle time ( NOOP ), for every benchmark ran, even with cache and pipeline on. This figure could have been improved if we had decided to implement a longer pipe with result forwarding. This demonstrates the trade off between design complexity and performance.

We also learned that it is important to have faster memory or cache. In our exchange sort benchmark, our simulator performed worse with the pipeline when cache was turned off. This was due to the fact that fetch must finish reading an instruction once started, even if it's going to get squashed. With slow memory and optimization such as a pipeline, a lot of work is wasted. Maybe some sort of prefetcher or predictor could help solve this, or a memory which allows changing of the address during reading.